Learning to Act in Decentralized Partially Observable MDPs

Jilles S. Dibangoye [1]   **Olivier Buffet** [2]

Abstract

We address a long-standing open problem of reinforcement learning in decentralized partially observable Markov decision processes. Previous attempts focussed on different forms of generalized policy iteration, which at best led to local optima. In this paper, we restrict attention to plans, which are simpler to store and update than policies. We derive, under certain conditions, the first near-optimal cooperative multi-agent reinforcement learning algorithm. To achieve significant scalability gains, we replace the greedy maximization by mixed-integer linear programming. Experiments show our approach can learn to act near-optimally in many finite domains from the literature.

1. Introduction

Decentralized partially observable Markov decision processes (Dec-POMDPs) emerged as the standard framework for sequential decision making by a team of collaborative agents (Bernstein et al., 2000). A key assumption of Dec-POMDPs is that agents can neither see the actual state of the system nor explicitly communicate their noisy observations with each other due to communication cost, latency or noise, hence providing a partial explanation of the double exponential growth at every control interval of the required memory in optimal algorithms (Hansen et al., 2004; Szer et al., 2005; Oliehoek et al., 2008; Amato et al., 2009; Oliehoek et al., 2013; Dibangoye et al., 2015; 2016). While planning methods for finite Dec-POMDPs made substantial progress in recent years, the formal treatment of the corresponding reinforcement learning problems received little attention so far. The literature of multi-agent reinforcement learning (MARL) can be divided into two main categories: *concurrent* and *team approaches* (Tan, 1998; Panait & Luke, 2005).

[1]Univ Lyon, INSA Lyon, INRIA, CITI, F-69621 Villeurbanne, France [2]INRIA / Université de Lorraine, Nancy, France. Correspondence to: Jilles S. Dibangoye <jilles.dibangoye@inria.fr>, Olivier Buffet <olivier.buffet@loria.fr>.

Perhaps the dominant paradigm in MARL is the concurrent approach, which involves multiple simultaneous learners: typically, each agent has its learning process. Self-interested learners, for example, determine their best-response behaviors considering their opponents are part of the environment, often resulting in local optima (Brown, 1951; Hu & Wellman, 1998; Littman, 1994). While concurrent learning can apply in Dec-POMDPs, a local optimum may lead to severely suboptimal performances (Peshkin et al., 2000; Zhang & Lesser, 2011; Kraemer & Banerjee, 2016; Nguyen et al., 2017). Also, methods of this family face two conceptual issues that limit their applicability. The primary concern is that of the co-adaptation dilemma, which arises when each attempt to modify an agent behavior can ruin learned behaviors of its teammates. Another major problem is that of the multi-agent credit assignment, that is, how to split the collective reward among independent learners.

Alternatively, the team approach involves a single learner acting on behalf of all agents to discover a collective solution (Salustowicz et al., 1998; Miconi, 2003). Interestingly, this approach circumvents the difficulties arising from both the co-adaptation and the multi-agent credit assignment. Coordinated agents, for example, simultaneously learn their control choices and the other agent strategies assuming instantaneous and free explicit communications (Guestrin et al., 2002; Kok & Vlassis, 2004). While methods of this family inherit from standard single-agent techniques, they need to circumvent two significant drawbacks: the explosion in the state space size; and the centralization of all learning resources in a single place. Recently, team algorithms ranging from Q-learning to policy-search have been introduced for finite Dec-POMDPs, but with no guaranteed global optimality (Wu et al., 2013; Liu et al., 2015; 2016; Kraemer & Banerjee, 2016). So, it seems one can either compute local optima with arbitrary bad performances or calculate optimal solutions but assuming noise-free, instantaneous and explicit communications.

A recent approach to optimally solving Dec-POMDPs suggests recasting them into occupancy-state MDPs (oMDPs) and then applying (PO)MDP solution methods (Dibangoye et al., 2013; 2014a;b; 2016). In these oMDPs, the states called occupancy states are distributions over hidden states and joint histories of the original problem, and actions called decision rules are mappings from joint histories to con-

trols (Nayyar et al., 2011; Oliehoek, 2013; Dibangoye et al., 2016). This approach achieves scalability gains by exploiting the piece-wise linearity and convexity of the optimal value function. Since this methodology was successfully applied for planning in Dec-POMDPs, it is natural to wonder which benefits it could bring to the corresponding MARL problem. Unfortunately, a straightforward application of standard RL methods to oMDPs will face three severe limitations. First, occupancy states are unknown, and hence must be estimated. Second, they lie in a continuum making tabular RL methods inapplicable. Finally, the greedy maximization is computationally demanding in decentralized stochastic control problems (Radner, 1962; Dibangoye et al., 2009; Kumar & Zilberstein, 2009; Oliehoek et al., 2010).

This paper extends the methodology of Dibangoye et al. to MARL, focussing on the three major issues that limit its applicability. Our primary result is the proof that, by restricting attention to plans instead of policies, a linear function over occupancy states and joint decision rules, which is simple to store and update, can capture the optimal performance for Dec-POMDPs. We further use plans instead of policies in a policy iteration algorithm, with the plan always being improved with respect to a linear function and a linear function always being driven toward the linear function for the plan. Under accurate estimation of the occupancy states, the resulting algorithm, called occupancy-state SARSA (oSARSA) (Rummery, G. A. and Niranjan, 1994), is guaranteed to converge with probability one to a near-optimal plan for any finite Dec-POMDP. To extend its applicability to higher-dimensional domains, oSARSA replaces the greedy (or soft) maximization by a mixed-integer linear program for finite settings. Altogether, we obtain a MARL algorithm that can apply to finite Dec-POMDPs. Experiments show our approach can learn to act near-optimally in many finite domains from the literature.

We organize the remainder of this paper as follows. Section 2 extends a recent planning theory, starting with a formal definition of finite Dec-POMDPs. We proceed with the introduction of a framework for centralized MARL in Dec-POMDPs along with solutions to the three limitations mentioned above in Section 3. We further present the resulting algorithm oSARSA along with convergence guarantees in Section 4. Finally, we conduct experiments in Section 5, demonstrating our approach learns to act optimally in many finite domains from the literature. Proofs are provided in the companion technical report (Dibangoye & Buffet, 2018).

2. Planning in Dec-POMDPs as oMDPs

2.1. Finite Dec-POMDPs

A finite Dec-POMDP is a tuple $M \doteq (n, X, \{U^i\}, \{Z^i\}, p, r, \ell, \gamma, b_0)$, where n denotes the

number of agents involved in the decentralized stochastic control process; X is a finite set of hidden world states, denoted x or y; U^i is a finite private control set of agent $i \in [\![1; n]\!]$, where $U = U^1 \times \cdots \times U^n$ specifies the set of controls $u = (u^1, \ldots, u^n)$; Z^i is a finite private observation set of agent i, where $Z = Z^1 \times \cdots \times Z^n$ specifies the set of observations $z = (z^1, \ldots, z^n)$; p describes a transition function with conditional density $p^{u,z}(x, y)$; r is a reward model with immediate reward $r(x, u)$, we assume rewards are two-side bounded, i.e., for some $c \in \mathbb{R}^+$, $\forall x \in X, u \in U \colon |r(x, u)| \leqslant c$; ℓ is the planning horizon; $\gamma \in [0, 1]$ denotes the discount factor; and b_0 is the initial belief state with density $b_0(x_0)$. We shall restrict attention to finite planning horizon $\ell < \infty$ since an infinite planning horizon solution is within a small scalar $\epsilon > 0$ of a finite horizon solution where $\ell = \lceil \log_\gamma((1 - \gamma)\epsilon/c) \rceil$.

Because we are interested in MARL, we assume an incomplete knowledge about M, *i.e.,* p and r are either unavailable or only through a generative model. Hence, the goal of solving M is to find a plan, *i.e.,* a tuple of individual decision rules, one for each agent and time step: $\rho \doteq (a^1_{0:\ell}, \ldots, a^n_{0:\ell})$. A tth individual decision rule $a^i_t \colon O^i_t \mapsto \mathcal{P}(U^i)$ of agent i prescribes private controls based on the whole information available to the agent up to the tth time step, *i.e.,* history of controls and observations $o^i_t = (u^i_{0:t-1}, z^i_{1:t})$, where $o^i_0 = \varnothing$ and $o^i_t \in O^i_t$. A tth joint decision rule, denoted $a_t \colon O_t \mapsto \mathcal{P}(U)$, can be specified as $a_t(u|o) \doteq \prod_{i=1}^n a^i_t(u^i|o^i)$, where $O_t \doteq O^1_t \times \cdots \times O^n_t$, $o^i \in O^i_t$ and $o \doteq (o^1, \ldots, o^n) \in O_t$. From control interval t onward, agents collectively receive discounted cumulative rewards, denoted by random variable $R_t \doteq \lambda_1 r_t + \cdots + \lambda_\ell r_\ell$, where λ_t denotes the time-step dependent weighting factors, often set to $\lambda_t = \gamma^t$ for discounted problems. For any control interval t, joint plans $a_{0:t}$ of interest are those that achieve the highest performance measure $J(a_{0:t}) \doteq \mathbb{E}^{a_{0:t}} \{R_0 \mid b_0\}$ starting at b_0, where $\mathbb{E}^{a_{0:t}}\{\cdot\}$ denotes the expectation with respect to the probability distribution over state-action pairs joint plan $a_{0:t}$ induces, in particular $J(\rho) \doteq J(a_{0:\ell-1})$ for $\rho \doteq a_{0:\ell-1}$. One can show that, in Dec-POMDPs, there always exists a deterministic plan that is as good as any stochastic plan (see Puterman, 1994, Lemma 4.3.1). Unfortunately, there is no direct way to apply the theory developed for Markov decision processes (Bellman, 1957; Puterman, 1994) to Dec-POMDPs, including: the Bellman optimality equation; or the policy improvement theorem. To overcome these limitations, we rely on a recent theory by Dibangoye et al. that recasts M into an MDP, thereby allowing knowledge transfer from the MDP setting to Dec-POMDPs.

2.2. Occupancy-State MDPs

To overcome the fact that agents can neither see the actual state of the system nor explicitly communicate their noisy observations with each other, Szer et al. (2005) and later

on Dibangoye et al. (2016) suggest formalizing M from the perspective of a centralized algorithm. A centralized algorithm acts on behalf of the agents by selecting a joint decision rule to be executed at each control interval based on all data available about the system, namely the information state. The information state at the end of control interval t, denoted $\iota_{t+1} \doteq (b_0, a_{0:t})$, is a sequence of joint decision rules the centralized algorithm selected starting at the initial belief state. Hence, the information state satisfies the following recursion: $\iota_0 \doteq (b_0)$ and $\iota_{t+1} \doteq (\iota_t, a_t)$ for all control interval t, resulting in an ever-growing sequence. To generalize the value from one information state to another one, Dibangoye et al. introduced the concept of occupancy states. The occupancy state at control interval t, denoted $s_t \doteq \mathbb{P}(x_t, o_t | \iota_t)$, is a distribution over hidden states and joint histories conditional on information state ι_t at control interval t. Interestingly, the occupancy state has many important properties. First, it is a sufficient statistic of the information state when estimating the (current and future) reward to be gained by executing a joint decision rule:

$$R(s_t, a_t) \doteq \sum_x \sum_o s_t(x, o) \sum_u a_t(u|o) \cdot r(x, u).$$

In addition, it describes a deterministic and fully observable Markov decision process, where the next occupancy state depends only on the current occupancy state and next joint decision rule, for all $y \in X, o \in O, u \in U, z \in Z$:

$$T(s_t, a_t) \doteq s_{t+1}$$
$$s_{t+1}(y, (o, u, z)) \doteq a_t(u|o) \sum_x s_t(x, o) \cdot p^{u,z}(x, y).$$

The process the occupancy states describe is known as the occupancy-state Markov decision process (*o*MDP), and denoted $M' \doteq (S, A, R, T, \ell, \gamma, s_0)$. Similarly to POMDPs, it was proven that Dec-POMDPs can be recasted into MDPs, called *o*MDPs, and an optimal solution of the resulting *o*MDP is also an optimal solution for the original Dec-POMDP (Dibangoye et al., 2016). M' is an ℓ-steps deterministic and continuous MDP with respect to M, where $S \doteq \cup_{t \in [\![0;\ell-1]\!]} S_t$ is the set of occupancy states up to control interval $\ell - 1$; $A \doteq \cup_{t \in [\![0;\ell-1]\!]} A_t$ is the set of joint decision rules up to control interval $\ell - 1$; R is the reward model; and T is the transition rule; s_0 is the initial occupancy state, which is essentially the initial belief in M; γ and ℓ are as in M. It is worth noticing that there is no need to construct explicitly M'; instead we use M (when available) as a generative model for the occupancy states $T(s_t, a_t)$ and rewards $R(s_t, a_t)$, for all control intervals t.

To better understand why we use plans instead of policies and how they relate, consider the MDP case. The solution of any finite MDP called a policy $\pi \colon S \mapsto A$ can be represented as a decision tree, where nodes are labelled with actions and arcs are labelled with states. Since an *o*MDP is also an MDP, its policies can also be represented as decision trees, except that actions are decision rules and states are occupancy states. In contrast to standard MDPs, *o*MDPs are deterministic. This means that only a single branch in the decision-tree representation—i.e., a sequence of actions—is necessary to act optimally in *o*MDPs. A single branch of a decision tree is called a plan. Hence policies are more general than plans, but in deterministic MDPs, both can be employed while achieving optimal performance (plans inducing an open-loop approach, and policies a closed-loop approach). We shall restrict attention to plans because they are more concise than policies. Below, we review a closed-loop approach based on the dynamic programming theory (Bellman, 1957).

For any finite M, the Bellman equation is written as follows: for all occupancy state $s_t \in S_t$, and some fixed policy π,

$$V_t^\pi(s_t) \doteq R(s_t, \pi(s_t)) + \lambda_1 V_{t+1}^\pi(T(s_t, \pi(s_t))) \quad (1)$$

with boundary condition $V_\ell^\pi(\cdot) \doteq 0$, describes the return of a particular occupancy state s_t when taking decision rule $a_t = \pi(s_t)$ prescribed by π. The equation for an optimal policy π^* is referred to as the Bellman optimality equation: for any control interval t, and occupancy state s_t,

$$V_t^*(s_t) \doteq \max_{a_t \in A} R(s_t, a_t) + \lambda_1 V_{t+1}^*(T(s_t, a_t)) \quad (2)$$

with boundary condition $V_\ell^*(\cdot) \doteq 0$. Unfortunately, occupancy states lie in a continuum, which makes exact dynamic programming methods infeasible. Interestingly, when optimized exactly, the value function solution of (2) along with the boundary condition is always piece-wise linear and convex in the occupancy-state space (Dibangoye et al., 2016).

Lemma 1. *For any arbitrary M', the solution $V_{0:\ell}^*$ of (2) is convex in the occupancy-state space. If we restrict attention to deterministic policies and finite M (and corresponding M'), the solution of (2) is piece-wise linear and convex in the occupancy-state space. Hence, the optimal value at any occupancy state s_t is as follows:*

$$V_t^*(s_t) \doteq \max_{\alpha_t \in \Gamma_t} \langle s_t, \alpha_t \rangle, \quad (3)$$

where $\langle s_t, \alpha_t \rangle$ is used to express the expectation of a linear function α_t (also called α-vector in the probability space defined by sample space $X \times O$, the σ-algebra $X \times O$ and the probability distribution s_t; and Γ_t is the set of all tth α-vectors.

Lemma 1 shows that for any arbitrary M and corresponding M', the solution of (2), represented by sets $\Gamma_{0:\ell}$, is convex in the occupancy-state space. Each α-vector defines the value function over a bounded region of the occupancy-state space. In addition, it is associated with a plan, defining the

optimal plan for a bounded region of the occupancy-state space. Sets $\Gamma_{0:\ell}$ are iteratively improved by adding a new α-vector that dominates current ones over certain regions of the occupancy-state space. The α-vector to be added is computed using point-based Bellman backup operator $\mathbb{H}$:

$$[\mathbb{H}\Gamma_{t+1}](s_t) = \arg\max_{\alpha_t^a \,:\, a \in A_t, \alpha_{t+1} \in \Gamma_{t+1}} \langle s_t, \alpha_t^a \rangle,$$

where $\alpha_t^a(x, o) \doteq \mathbb{E}\{r(x, u) + \lambda_1 \alpha_{t+1}(y, (o, u, z)) | a\}$, for each hidden state $x \in X$, and joint history $o \in O$. To keep the number of α-vectors manageable, one can prune those that are dominated over the entire occupancy-state space. All in all, the oMDP reformulation permits us to solve finite M by means of M' using near-optimal planning methods leveraging on the special structure of the optimal value function (Shani et al., 2013). This methodology results in the current state-of-the-art algorithm to optimally solving finite Dec-POMDPs (Dibangoye et al., 2016). So it seems natural to wonder if the same methodology can also succeed when applied to the corresponding reinforcement-learning problem. In other words, how can a centralized algorithm learn to coordinate a team of agents with possibly contradicting perceptual information?

3. Learning in Dec-POMDPs as oMDPs

Using the oMDP reformulation, a natural approach to achieve centralized RL for decentralized stochastic control suggests applying exact RL methods. In the Q-learning algorithm (Watkins & Dayan, 1992), for example, one would learn directly the Q-value function when following a fixed policy π: for any control interval $t \in [\![0; \ell - 1]\!]$,

$$Q_t^\pi(s_t, a_t) \doteq R(s_t, a_t) + \lambda_1 V_{t+1}^\pi(T(s_t, a_t)) \quad (4)$$

with boundary condition $Q_\ell^\pi(\cdot, \cdot) = 0$. The policy improvement theorem provides a procedure to change a sub-optimal policy π into an improved one $\bar{\pi}$ (Howard, 1960): for any control interval $t \in [\![0; \ell - 1]\!]$,

$$\bar{\pi}(s_t) \doteq \arg\max_{a_t \in A_t} Q_t^\pi(s_t, a_t). \quad (5)$$

Unfortunately, this approach has three severe limitations. First, the occupancy states are unknown and must be estimated. Second, even if we assume a complete knowledge of the occupancy states, they lie in a continuum, which precludes exact RL methods to accurately predict α-vectors even in the limit of infinite time and data. Finally, the greedy maximization required to improve the value function proved to be NP-hard in finite settings (Radner, 1962; Dibangoye et al., 2009; Kumar & Zilberstein, 2009; Oliehoek et al., 2010).

3.1. Addressing Estimation Issues

Although mappings T and R in M' are unknown to either agents or a centralized algorithm, one can instead estimate

on the fly both $T(s_0, a_{0:t-1})$ and $R(T(s_0, a_{0:t-1}), a_t)$ for some fixed plan $\rho \doteq a_{0:\ell-1}$ through successive interactions of agents with the environment. To this end, we shall distinguish between two settings. The first one assumes a generative model is available during the centralized learning phase, e.g. a black box simulator; and the second does not. In both cases, we build on the concept of replay pool (Mnih et al., 2015), except that we extend it from stationary single-agent domains to non-stationary multi-agent domains.

If a generative model is available during the learning phase, then a Monte Carlo method can approximate $T(s_0, a_{0:t-1})$ and $R(T(s_0, a_{0:t-1}), a_t)$ arbitrary closely. To this end, the generative model allows the agents to sample experiences generated from M. An ℓ-steps experience is a 4-tuple $\xi \doteq (x_{0:\ell-1}, u_{0:\ell-1}, r_{0:\ell-1}, z_{1:\ell})$, where $x_{0:\ell-1}$ are sampled hidden states, $u_{0:\ell-1}$ are controls made, $r_{0:\ell-1}$ are reward signals drawn from the reward model, and $z_{1:\ell}$ are the resulting observations, drawn from the dynamics model. If we let $\mathcal{D}^\rho \doteq \{\xi^{[i]}\}_{i \in [\![1:K]\!]}$ be the replay pool of K i.i.d random samples created through successive interactions with the generative model, then empirical occupancy state $\hat{s}_t \approx T(s_0, a_{0:t-1})$ and reward $\hat{R}_t \approx R(T(s_0, a_{0:t-1}), a_t)$ corresponding to the current $\mathcal{D}^\rho$ are given by: for any control interval $t \in [\![0 : \ell - 1]\!]$,

$$\hat{s}_t(x, o) \doteq \tfrac{1}{K} \sum_{i=1}^{K} \delta_x(x_t^{[i]}) \cdot \delta_o(u_{0:t}^{[i]}, z_{1:t}^{[i]}) \quad (6)$$

$$\text{and} \quad \hat{R}_t \doteq \tfrac{1}{K} \sum_{i=1}^{K} r_t^{[i]}, \quad (7)$$

where $\delta_x(\cdot)$ and $\delta_o(\cdot)$ denote the delta-Dirac mass located in hidden state and joint history pair, respectively. By the law of large numbers the sequence of averages of these estimates converges to their expected values, and the standard-deviation of its error falls as $1/\sqrt{K}$ (Sutton & Barto, 1998, chapter 5). The error introduced by Monte Carlo when estimating $T(\hat{s}_{t-1}, a_{t-1})$ instead of $T(s_0, a_{0:t-1})$ is upper bounded by $2\ell/\sqrt{K}$. The proof follows from the performance guarantee of the policy-search algorithm by Bagnell et al. (2004). Hence, to ensure the learned value function is within $\epsilon > 0$ of the optimal one, one should set the replay-pool size to $K = \Theta(4\frac{\ell^2}{\epsilon^2})$.

When no generative model is available, the best we can do is to store samples agents collected during the learning phase into replay pools $\mathcal{D}^\rho$, one experience for each episode within the limit size of K. We maintain only the K recent experiences, and may discard[1] hidden states since they are unnecessary for the updates of future replay pools and the performance measure. The rationale behind this approach is that it achieves the same performances as a Monte Carlo method for the task of approximating $T(s_0, a_{0:t-1})$ and $R(T(s_0, a_{0:t-1}), a_t)$ given a fixed plan $\rho \doteq a_{0:\ell-1}$. In fact, if we let $\mathcal{D}^\rho$ be a replay pool of K i.i.d. samples

---

[1] Note that one should keep hidden states when available since they often speed up the convergence.

generated according to ρ, the empirical occupancy state $\hat{s}_t \approx T(s_0, a_{0:t-1})$ and reward $\hat{R}_t \approx R(T(s_0, a_{0:t-1}), a_t)$ corresponding to $\mathcal{D}^\rho$ are given by (6) and (7), respectively. One can further show this approach preserves performance guarantees similar to those obtained when using a generative model.

3.2. Addressing Prediction Issues

The key issue with large spaces of occupancy states and decision rules is that of generalization, that is, how experiences with a limited subset of occupancy states and decision rules can produce a good approximation over a much larger space. Fortunately, a fundamental property of oMDPs is the convexity of the optimal value function over the occupancy-state space, see Lemma 1. Building on this property, we demonstrate a simple yet important preliminary result before stating the main result of this section.

Lemma 2. *For any arbitrary M' (resp. M), the optimal Q-value function is the upper envelope of sets $\Omega^*_{0:\ell}$ of α-vectors over occupancy states and joint decision rules: for any control interval t, $Q^*_t(s_t, a_t) = \max_{q_t \in \Omega^*_t} \langle s_\tau \odot a_\tau, q_t \rangle$, where $q_t \in \Omega^*_t$ are appropriate α-vectors, and $s_\tau \odot a_\tau$ denotes the Hadamard product[2].*

Lemma 2 generalizes the convexity property demonstrated in Lemma 1 from optimal value functions over occupancy states to optimal value functions over occupancy states and decision rules. As a consequence, finite sets $\Omega^*_{0:\ell-1}$ of α-vectors can produce solutions arbitrarily close to the optimal Q-value function $Q^*_{0:\ell-1}$. Though Q-value function $Q^*_{0:\ell-1}$ generalizes from a pair of occupancy state and decision rule to another one, storing and updating a convex hull is non trivial. Instead of learning the optimal Q-value function over all occupancy states and decision rules, we explore a simpler yet tractable alternative, which will prove sufficient to preserve ability to eventually find an optimal plan starting at initial occupancy state s_0.

Theorem 1. *For any arbitrary M' (resp. M), the Q-value function $Q^{\rho^*}_{0:\ell-1}$ under an optimal plan $\rho^* \doteq a^*_{0:\ell-1}$ starting at initial occupancy state s_0 is linear in occupancy states and decision rules: $Q^{\rho^*}_t(s_t, a_t) = \langle s_t \odot a_t, q^{\rho^*}_t \rangle$ where $q^{\rho^*}_t \doteq \arg\max_{q_t \in \Omega^*_t} \langle T(s_0, a^*_{0:t-1}) \odot a^*_t, q_t \rangle$.*

Theorem 1 proves that the Q-function for a given optimal joint plan achieves performance at the initial occupancy state s_0 as good as the Q-value function for an optimal joint policy. Standard policy iteration algorithms search for an optimal joint policy, which requires a finite set of α-vectors to approximate V^*/Q^*, hence the resulting PWLC approximator is tight almost everywhere. Building upon Theorem 1, we search for an optimal ρ, which requires

only a single α-vector to approximate V^ρ/Q^ρ, thus the resulting linear approximator is loose everywhere except in the neighborhood of a few points. The former approach may require less iterations before convergence to an optimal joint policy, but the computational cost of each iteration shall increase with the number of α-vectors maintained. The latter approach may require much more iterations, but all iteration shares the same computational cost.

3.3. Addressing Plan Improvement Issues

A fundamental theorem in many RL algorithms is the policy improvement theorem, which helps improving policies over time until convergence. This section introduces a procedure to improve a plan starting with a sub-optimal one.

Suppose we have determined the value function $V^\rho_{0:\ell-1}$ for any arbitrary $\rho \doteq a_{0:\ell-1}$. For some control interval $t \in [\![0; \ell-1]\!]$, we would like to know whether or not we should change decision rules $a_{0:t}$ to choose $\bar{a}_{0:t} \neq a_{0:t}$. We know how good it is to follow the current plan from control interval t onward—that is V^ρ_t—but would it be better or worse to change to the new plan? One way to answer this question is to consider selecting $\bar{a}_{0:t}$ at control interval t and thereafter following decision rules $a_{t+1:\ell-1}$ of the existing ρ. The value of the resulting joint plan is given by $J(\bar{a}_{0:t-1}) + \lambda_1 V^\rho_{t+1}(T(s_0, \bar{a}_{0:t}))$. The key criterion is whether this quantity is greater or less than $J(\rho)$. Next, we state the plan improvement theorem for oMDPs.

Theorem 2. *Let $\rho \doteq a_{0:\ell-1}$ and $\bar{\rho} \doteq \bar{a}_{0:\ell-1}$ be any pair of plans and $J_{0:\ell}$ be a sequence of α-vectors such that, for all t, $J_t(x_t, o_t) \doteq \mathbb{E}\{\alpha_0 r_0 + \ldots + \alpha_t r_t | b_0, x_t, o_t, a_{0:t-1}\}$. Let $\bar{s}_t \doteq T(s_0, \bar{a}_{0:t-1})$ and $s_t \doteq T(s_0, a_{0:t-1})$ be occupancy states at any control interval $t \in [\![0; \ell-1]\!]$ under $\bar{\rho}$ and ρ, respectively. Then, $\langle \bar{a}_{0:t^*-1}, a_{t^*:\ell-1} \rangle$ such that $t^* = \arg\max_{t \in [\![0;\ell-1]\!]} \langle \bar{s}_t - s_t, J_t - \lambda_1 V^\rho_t \rangle$ is as good as, or better than, ρ.*

Proof. The proof follows from the difference between the performance measure of $\rho \doteq a_{0:\ell-1}$ and $\bar{\rho} \doteq \bar{a}_{0:\ell-1}$. Let $\varsigma_t(\bar{\rho}, \rho)$ be the advantage of taking plan $\langle \bar{a}_{0:t-1}, a_{t:\ell-1} \rangle$ instead of ρ: for any control interval $t \in [\![0; \ell-1]\!]$,

$$\varsigma_t(\bar{\rho}, \rho) = J(\bar{a}_{0:t-1}) + \lambda_1 V^\rho_t(T(s_0, \bar{a}_{0:t-1})) - J(\rho)$$
$$= J(\bar{a}_{0:t-1}) - J(a_{0:t-1}) + \lambda_1(V^\rho_t(\bar{s}_t) - V^\rho_t(s_t))$$
$$= \langle \bar{s}_t - s_t, J_t - \lambda_1 V^\rho_t \rangle.$$

If we let $t^* \doteq \arg\max_{t=0,1,\ldots,\ell-1} \varsigma_t(\bar{\rho}, \rho)$, then plan $\langle \bar{a}_{0:t^*-1}, a_{t^*:\ell-1} \rangle$ achieves the highest advantage among plan set $\{\langle \bar{a}_{0:t-1}, a_{t:\ell-1} \rangle\}_{t \in [\![0;\ell-1]\!]}$ constructed based on $\bar{\rho}$. If $t^* = 0$, then $\langle \bar{a}_{0:t^*-1}, a_{t^*:\ell-1} \rangle = \rho$, and no improved plans were found from plan set generated from $\bar{\rho}$. Otherwise, new $\langle \bar{a}_{0:t^*-1}, a_{t^*:\ell-1} \rangle$ must be better than ρ. $\square$

Theorem 2 plays the same role in the plan space as does

---

[2] $\forall(x, o, u) \colon [s_\tau \odot a_\tau](x, o, u) \doteq s_\tau(x, o) \cdot a_\tau(u|o)$.

the policy improvement theorem in the policy space. More precisely, after sampling a plan, i.e., a sequence of decision rules, it tells us which of these decision rules will improve the current plan. More specifically, it shows how, given $\rho \doteq a_{0:\ell-1}$ and α-vector $q_{0:\ell-1}^{\rho}$, we can easily evaluate a change in ρ at any control interval to a particular (possibly improved) plan. To ease exploration towards promising plans, we investigate the ϵ-greedy maximization (or soft-maximization). At each control interval t and occupancy state s_t, it randomly selects $\hat{a}_t$ with probability ϵ; otherwise, it greedily selects $\hat{a}_t$ w.r.t. the current Q-value function Q_t^{ρ}:

$$\hat{a}_t \doteq \arg\max_{a_t : a_t^1 \in A_t^1, \dots, a_t^n \in A_t^n} Q_t^{\rho}(s_t, a_t),$$

where $\hat{\rho} \doteq \hat{a}_{0:\ell-1}$. Unfortunately, this operation proved to be NP-hard for finite M (Radner, 1962; Dibangoye et al., 2009; Kumar & Zilberstein, 2009; Oliehoek et al., 2010). Searching for the best decision rule requires enumerating all of them, which is not possible in large planning horizons. Instead, we present a mixed-integer linear programming (MILP) method, which successfully performs the greedy maximization for finite M. Though MILP is NP-hard in the worst case, the solution of its LP relaxation, which is in P, is often integral in our experiments. In other words, the solution of the LP relaxation is already a solution of the MILP. A similar observation was done before by Mac-Dermed & Isbell. Mixed-Integer Linear Program 1 builds on (MacDermed & Isbell, 2013), which introduced an integer program for the greedy maximization in finite M. We also exploit the occupancy state estimation, in which $\hat{s}_t$ replaces s_t, and the current α-vector q_t^{ρ}.

Mixed-Integer Linear Program 1 (For finite M).

$$\underset{a_t, a_t^1, \dots, a_t^n}{Maximize} \sum_x \sum_o \hat{s}_t(x, o) \sum_u a_t(u|o) \cdot q_t^{\rho}(o, u) \quad (8)$$

$$s.t.: \sum_{u^j} a_t(u^j, u^i|o) = a_t^i(u^i|o^i), \ \forall i, u^i, o \quad (9)$$

$$\sum_u a_t(u|o) = 1, \ \forall o \quad (10)$$

where $\{a_t(u|o)\}$ and $\{a_t^i(u^i|o^i)\}$ are positive and boolean variables, respectively.

Mixed-Integer Linear program 1 optimizes positive variables $\{a_t(u|o)\}_{u \in U, o \in O_t}$, one positive variable for each control-history pair. More precisely, each variable represents the probability $a_t(u|o)$ of control u being taken given that agents experienced joint history o. Constraints must be imposed on these variables to ensure they form proper probability distributions (10), and that they result from the product of independent probability distributions (9), one independent probability distribution for each agent. In order to make the description of the conditional independence,

$$a_t(u|o) = a_t^1(u^1|o^1) \times \dots \times a_t^n(u^n|o^n), \quad (11)$$

we use additional variables $\{a_t^i(u^i|o^i)\}_{i \in [\![1;n]\!], u^i \in U^i, o^i \in O_t^i}$. Marginalizing out both sides of (11) over all control-history

pairs of all agents except agent i, denoted $-i$, leads to (9). That is not sufficient to ensure conditional independence in general. If we further constrain $\{a_t^i(u^i|o^i)\}$ to be boolean, then system of equations (9) implies (11). Given (9) and (10), agent variables $\{a_t^i(u^i|o^i)\}_{u^i \in U^i, o^i \in O_t^i}$ describe a proper probability distribution, so we omit corresponding constraints. Our greedy maximization approach is fundamentally different from previous ones, including the integer program by (MacDermed & Isbell, 2013) and the constraint optimization program by (Kumar & Zilberstein, 2009; Dibangoye et al., 2016). First, while previous approaches made use of boolean variables, we use both positive and boolean variables instead. Next, prior approaches optimize a value function represented as a convex hull; we optimize an α-vector instead.

4. The oSARSA Algorithm

This section presents the oSARSA algorithm with tabular representations and function approximations (using either linear functions or deep neural networks) along with convergence guarantees. oSARSA algorithms are specializations of Policy Iteration, except that we use plans instead of policies. For the sake of conciseness, we describe a generic algorithm, which can fit to either tabular or approximate representations.

In Dec-POMDPs, the goal of oSARSA is to learn $q_{0:\ell-1}^*$, a sequence of α-vectors of an optimal plan ρ^*. In particular, we must estimate $q_t(x, o, u)$ for the current plan ρ and for all reachable state x, joint history o, control u, and any control interval t. At the same time, the algorithm changes ρ towards improved plans according to the plan improvement theorem. The improved plans are constructed by exploring the occupancy-state space according to ϵ-greedy plans (see Section 3.3). To provide good estimations, we store all experiences in data set $\mathcal{D}^{\rho}$, from which we estimate the occupancy states and returns under ρ for any control interval (see Section 3.1). Upon estimating occupancy state $\hat{s}$ and selecting joint decision rule a, we update parametrized α-vector q_t with parameter θ_t using q_{t+1}, $\mathcal{D}^{\rho}$ and a_{t+1} by means of temporal difference learning: for all (x, o, u),

$$\theta_t^{[\tau+1]} \doteq \theta_t^{[\tau]} + \beta_\tau \mathbb{E}_{\hat{s}, a, \mathcal{D}, a_{t+1}} \{\delta_t \nabla q_t^{[\tau]}(x, o, u)\} \quad (12)$$

$$\delta_t = r + \lambda_1 q_{t+1}^{[\tau]}(y, o', u') - q_t^{[\tau]}(x, o, u),$$

where β_τ is a step size, and quantity $\nabla q_t(x, o, u)$ denotes the gradient of q_t at (x, o, u) w.r.t. some parameter θ_t. Using tabular representations (e.g., finite/small M), $\theta_t = q_t$ and thus $\nabla q_t(x, o, u)$ is a unit vector $e_{x,o,u}$ whose value at (x, o, u) is one and zero otherwise. Using linear function approximations (e.g., continuous/large M), $q_t(x, o, u) \doteq \phi_t(x, o, u)^\top \theta_t$, where $\nabla q_t(x, o, u) = \phi_t(x, o, u)$ is the feature vector at (x, o, u). Algorithm 1 shows the pseudocode of oSARSA.

Algorithm 1 The oSARSA Algorithm

    Initialize $\bar{g} = -\infty$, $\bar{\rho}$ and $q_{0:\ell-1}$ arbitrary, and $\mathcal{D}^{\bar{\rho}}$.
    while $q_{0:\ell-1}$ has not converged **do**
        Select ϵ-greedily ρ w.r.t. $q_{0:\ell-1}$ and $\mathcal{D}^{\bar{\rho}}$.
        Compose $\mathcal{D}^{\rho}$ with N trajectories $\{\xi^{[\tau]}\}_{\tau=1}^{N}$.
        Estimate (g, ς) from $[\sum_{t=0}^{\ell-1} \hat{R}_t | \mathcal{D}^{\rho}, \hat{s}_0 = s_0]$.
        If $g - \varsigma \geqslant \bar{g}$ **then** $(\bar{\rho}, \bar{g}, \mathcal{D}^{\bar{\rho}}) = (\rho, g + \varsigma, \mathcal{D}^{\rho})$.
        Update α-vectors $q_{0:\ell-1}$ as described in (12).
    end while

To establish the convergence of oSARSA, we introduce the following assumptions.

Theorem 3. *Consider assumptions: (1) The stepsizes $\{\beta_\tau\}_{\tau=1,2,\ldots}$ satisfy Robbins & Monro's conditions; (2) The occupancy states $\hat{s}_{0:\ell-1}$ and immediate returns $\hat{R}_{0:\ell-1}$ are accurately estimated; and (3) Every pair of reachable occupancy state and joint decision rule is visited infinitely often. Under these assumptions, the sequence $q_{0:\ell-1}^{[\tau]}$ generated by oSARSA converges with probability 1 to $q_{0:\ell-1}^{*}$.*

Proof. Under these assumptions, we define $\mathbb{H}^{\rho}$ that maps a sequence of α-vectors $q_{0:\ell-1}$ to a new sequence of α-vectors $\mathbb{H}^{\rho} q_{0:\ell-1}$ according to the formula: for all hidden state x, joint history o and control u, at control interval t,

$$(\mathbb{H}^{\rho} q_{0:\ell-1})(x, o, u) = r(x, u) + \lambda_1 \mathbb{E}\{v_{t+1}(y, o \oplus (u, z))\},$$

where $v_t(x, o) \doteq q_t(x, o, \rho(o))$ and $\rho(o)$ is the control prescribed by ρ at joint history o. Then, the plan evaluation step of the oSARSA algorithm is of the form

$$q_t^{[\tau+1]}(x, o, u) = (1 - \beta_t) q_t^{[\tau]}(x, o, u) + \beta_t \kappa_t^{[\tau]}(x, o, u),$$
$$\kappa_t^{[\tau]}(x, o, u) = (\mathbb{H}^{\rho} q_{0:\ell-1}^{[\tau]})(x, o, u) + w_t(x, o, u),$$

where $w_t(x, o, u) = r(x, u) + \lambda_1 v_{t+1}^{[\tau]}(y, o \oplus (u, z)) - (\mathbb{H}^{\rho} q_{0:\ell-1}^{[\tau]})(x, o, u)$ is a zero mean noise term. Using this temporal-difference update-rule, see (12), we converge with probability 1 to $q_{0:\ell-1}^{\rho}$. It now remains to be verified that the plan improvement step of the oSARSA algorithm changes the current plan for an improved one. Initially, $\bar{g}$ is arbitrarily bad, so any new plan is an improved one. Then, $\bar{g} = J(\rho)$ for the current best plan ρ since occupancy state and return are accurately estimated. Hence, when ever $g \geqslant \bar{g}$, we know that the new plan $\bar{\rho}$ yields a performance measure $J(\bar{\rho})$ superior to $J(\rho)$, thus $\bar{\rho}$ improves ρ. We conclude the proof noticing that in finite M, the number of deterministic plans is finite. As a consequence, by visiting infinitely often every pair of occupancy state and decision rule we are guaranteed to visit all deterministic plans, hence an optimal one. $\square$

It is now important to observe that we meet assumption (2) in Theorem 3 only when M is available. Otherwise, we rely on confidence bounds $[g - \varsigma, g + \varsigma]$, e.g. *Hoeffding's inequality*, on estimate $g \approx J(\rho)$. In particular, we use lower-bounds $g - \varsigma$ on sample means instead of the sample means g themselves, to limit situations where g is overestimated. Small data sets often lead to suboptimal solutions, but as the number of experiences in data set $\mathcal{D}^{\rho}$ increases, sample means and corresponding lower bounds get close to the mean, i.e., ς tends to 0. We require an accurate estimation of a plan's performance to know for sure its performance is above that of any other plans we may encounter. However, an estimation of a sample plan's performance can be refined over time until it becomes accurate. For example, if the algorithm samples a promising plan—i.e., its confidence bounds suggest it might achieve a better performance than that of the current plan—the algorithm can progressively refine it until it becomes accurate. Of course, having a good initial estimate can significantly speed up the convergence. In the extreme case, the initial estimate is the true value. It is also worth noticing that the memory complexity of the oSARSA algorithms is linear with the size of an α-vector, *i.e.,* $\Theta(|\mathcal{D}^{\rho}|)$; and its time complexity is linear with the episodes.

5. Experiments

We ran the oSARSA algorithm on a Mac OSX machine with 3.8GHz Core i5 and 8GB of available RAM. We solved the MILPs using ILOG CPLEX Optimization Studio. We define features to use sequences of K last joint observations instead of joint histories, hence the dimension of the parameter vector θ is $|X|(|U||Z|)^K$ for finite M.

We evaluate our algorithm on multiple 2-agent benchmarks from the literature all available at `masplan.org`: Mabc, Recycling, Gridsmall, Grid3x3corners, Boxpushing, and Tiger. These are the largest and the most challenging benchmarks from the Dec-POMDP literature. For each of them, we compare our algorithm to the state-of-the-art algorithms based on either a complete or a generative model: FB-HSVI (Dibangoye et al., 2016), RLaR (Kraemer & Banerjee, 2016), and MCEM (Wu et al., 2013). We also reported results of the state-of-the-art *model-free* solver: (distributed) REINFORCE (Peshkin et al., 2000). For REINFORCE and oSARSA, we used hyper-parameters ϵ and β ranging from 1 to 10^{-3} with a decaying factor of 10^4, sample size $|\mathcal{D}| = 10^4$. We use maximum episodes and time limit 10^5 and 5 hours, respectively, as our stopping criteria.

Surprisingly, REINFORCE performs very well on domains that consist of weakly coupled agents, *see* Figure 1. However, for domains with strongly coupled agents, *e.g.,* Tiger or BoxPushing, it often gets stuck at some local optima. In contrast, oSARSA converges to near-optimal solutions when enough resources are available over all domains, *see* Figure 1 and Table 1. Regarding the most challenging benchmarks, which require more resources, oSARSA stops before

the convergence to a near-optimal solution; yet, it often outperforms the other RL algorithms. RLaR can achieve near-optimal result for small domains and short planning horizon ($\ell \leqslant 5$), assuming there exists a unique optimal plan. As for MCEM, it can solve infinite horizon problems, but similarly to REINFORCE may get stuck in local optima; this is essentially as they both use a form of gradient descent in a parametrized policy space.

T	RLaR	MCEM	REINFORCE	**oSARSA**	FB-HSVI
Tiger ($|X| = 2, |Z| = 4, |U| = 9, K = 3$)					
3	5.19	N.A.	5.0	5.19	5.19
4	4.46	N.A.	4.6	4.80	4.80
5	6.65	N.A.	2.2	6.99	7.02
6	–	N.A.	0.3	2.34	10.38
7	–	N.A.	-1.7	2.25	9.99
∞	N.A.	-10	-19.9	-0.2	13.44
Grid3x3corners ($|X| = 81, |Z| = 81, |U| = 25, K = 1$)					
6	–	N.A.	1.46	1.49	1.49
7	–	N.A.	2.17	2.19	2.19
8	–	N.A.	2.96	2.95	2.96
9	–	N.A.	3.80	3.80	3.80
10	–	N.A.	4.66	4.69	4.68
Boxpushing ($|X| = 100, |Z| = 16, |U| = 25, K = 1$)					
3	66.08	N.A.	17.6	65.27	66.08
4	98.59	N.A.	18.1	98.16	98.59
5	–	N.A.	35.2	107.64	107.72
6	–	N.A.	36.4	120.26	120.67
7	–	N.A.	36.4	155.21	156.42
8	–	N.A.	52.9	186.04	191.22
9	–	N.A.	54.5	206.75	210.27
10	–	N.A.	54.7	218.39	223.74
∞	N.A.	59.1	58.9	144.57	224.43

Table 1. Comparing $V^\rho(s_0)$ of all solvers when available, where the "–" sign mean "out of memory" and/or "out of time".


Figure 1. Comparing $V^\rho(s_0)$ of solvers with $\ell = \infty$ and $\lambda = 0.9$, where x-axis denotes the number of episodes during training.

6. Discussion

This paper extends a recent but growing (deep) MARL paradigm (Szer et al., 2005; Dibangoye et al., 2016; Kraemer & Banerjee, 2016; Mordatch & Abbeel, 2017; Foerster et al., 2017), namely RL for decentralized control, from model-based to model-free settings. This paradigm allows a centralized algorithm to learn on behalf of all agents how to select an optimal joint decision rule to be executed at each control interval based on all data available about the system during a learning phase, while still preserving ability for each agent to act based solely on its private histories at the execution phase. In particular, we introduced tabular and approximate oSARSA algorithms, which demonstrated promising results often outperforming state-of-the-art MARL approaches for Dec-POMDPs. To do so, oSARSA learns a value function that maps pairs of occupancy state and joint decision rule to reals. To ease the generalization in such high-dimensional continuous spaces, we restrict attention to plans rather than policies, which in turn restricts value functions of interest to linear functions. To speed up the greedy maximization, we used a MILP for finite settings—we shall use a gradient approach instead of a MILP for continuous settings in future works. Finally, we present a proof of optimality for a MARL algorithm when the estimation error is neglected. We shall investigate an approach to relax this somewhat restrictive assumption, perhaps within the probably approximately correct learning framework.

The RL for decentralized control paradigm is significantly different from the standard RL paradigm, in which agents have the same amount of information during both the learning and the execution phases. Another major difference lies in the fact that learned value functions in standard (deep) RL algorithms are mapping from histories (or states) to reals. In contrast, oSARSA learns a value function that maps occupancy-state/decision-rule pairs to reals—spaces of occupancy states and joint decision rules are multiple orders of magnitude larger than history or state spaces. As a consequence, standard (MA)RL methods, *e.g.* REINFORCE and MCEM, and recent actor-critic methods (Bono et al., 2018), may converge towards a local optimum faster than oSARSA, but the latter often converges towards a near-optimal solution. oSARSA uses occupancy states instead of joint histories mainly because occupancy states are (so far minimal) sufficient statistics for optimal decision-making in Dec-POMDPs—using joint histories instead of occupancy states may lead to suboptimal solutions except in quite restrictive settings. For example, RLaR learns value functions mapping history/action pairs to reals, but convergence towards an optimal solution is guaranteed only for domains that admit a unique optimal joint plan—which essentially restricts to POMDPs (Kraemer & Banerjee, 2016).

References

Amato, C., Dibangoye, J. S., and Zilberstein, S. Incremental Policy Generation for Finite-Horizon DEC-POMDPs. In *ICAPS*, 2009.

Bagnell, J. A., Kakade, S. M., Schneider, J. G., and Ng, A. Y. Policy Search by Dynamic Programming. In *NIPS*. 2004.

Bellman, R. E. *Dynamic Programming*. 1957.

Bernstein, D. S., Zilberstein, S., and Immerman, N. The Complexity of Decentralized Control of Markov Decision Processes. In *UAI*, 2000.

Bono, G., Dibangoye, J. S., Matignon, L., Pereyron, F., and Simonin, O. On the Study of Cooperative Multi-Agent Policy Gradient. Research Report RR-9188, INSA Lyon ; INRIA, June 2018. URL https://hal.inria.fr/hal-01821677.

Brown, G. W. Iterative Solutions of Games by Fictitious Play. In *Activity Analysis of Production and Allocation*. 1951.

Dibangoye, J. S. and Buffet, O. Learning to Act in Decentralized Partially Observable MDPs. Research report, INRIA Grenoble - Rhone-Alpes - CHROMA Team ; INRIA Nancy, équipe LARSEN, June 2018. URL https://hal.inria.fr/hal-01809897.

Dibangoye, J. S., Mouaddib, A.-I., and Chaib-draa, B. Point-based incremental pruning heuristic for solving finite-horizon DEC-POMDPs. In *AAMAS*, pp. 569–576, 2009.

Dibangoye, J. S., Amato, C., Buffet, O., and Charpillet, F. Optimally Solving Dec-POMDPs As Continuous-state MDPs. In *IJCAI*, pp. 90–96, 2013.

Dibangoye, J. S., Amato, C., Buffet, O., and Charpillet, F. Optimally solving Dec-POMDPs as Continuous-State MDPs: Theory and Algorithms. Research Report RR-8517, INRIA, April 2014a.

Dibangoye, J. S., Buffet, O., and Charpillet, F. Error-Bounded Approximations for Infinite-Horizon Discounted Decentralized POMDPs. In *ECML*, pp. 338–353, 2014b.

Dibangoye, J. S., Buffet, O., and Simonin, O. Structural Results for Cooperative Decentralized Control Models. In *IJCAI*, pp. 46–52, 2015.

Dibangoye, J. S., Amato, C., Buffet, O., and Charpillet, F. Optimally Solving Dec-POMDPs as Continuous-State MDPs. *Journal of AI Research*, 55, 2016.

Foerster, J. N., Farquhar, G., Afouras, T., Nardelli, N., and Whiteson, S. Counterfactual Multi-Agent Policy Gradients. *CoRR*, 2017.

Guestrin, C., Lagoudakis, M. G., and Parr, R. Coordinated Reinforcement Learning. In *ICML*, San Francisco, CA, USA, 2002.

Hansen, E. A., Bernstein, D. S., and Zilberstein, S. Dynamic Programming for Partially Observable Stochastic Games. In *AAAI*, 2004.

Howard, R. A. *Dynamic Programming and Markov Processes*. 1960.

Hu, J. and Wellman, M. P. Multiagent Reinforcement Learning: Theoretical Framework and an Algorithm. In *ICML*, San Francisco, CA, USA, 1998.

Kok, J. R. and Vlassis, N. Sparse Cooperative Q-learning. In *ICML*, New York, NY, USA, 2004.

Kraemer, L. and Banerjee, B. Multi-agent reinforcement learning as a rehearsal for decentralized planning. *Neurocomputing*, 190, 2016.

Kumar, A. and Zilberstein, S. Constraint-based dynamic programming for decentralized POMDPs with structured interactions. In *AAMAS*, 2009.

Littman, M. L. Markov games as a framework for multi-agent reinforcement learning. In *ICML*, 1994.

Liu, M., Amato, C., Liao, X., Carin, L., and How, J. P. Stick-breaking policy learning in Dec-POMDPs. In *IJCAI*. AAAI, 2015.

Liu, M., Amato, C., Anesta, E. P., Griffith, J. D., and How, J. P. Learning for Decentralized Control of Multiagent Systems in Large, Partially-Observable Stochastic Environments. In *AAAI*, 2016.

MacDermed, L. C. and Isbell, C. Point Based Value Iteration with Optimal Belief Compression for Dec-POMDPs. In *NIPS*, 2013.

Miconi, T. When Evolving Populations is Better Than Coevolving Individuals: The Blind Mice Problem. In *IJCAI*, 2003.

Mnih, V., Kavukcuoglu, K., Silver, D., Rusu, A. A., Veness, J., Bellemare, M. G., Graves, A., Riedmiller, M., Fidjeland, A. K., Ostrovski, G., Petersen, S., Beattie, C., Sadik, A., Antonoglou, I., King, H., Kumaran, D., Wierstra, D., Legg, S., and Hassabis, D. Human-level control through deep reinforcement learning. *Nature*, 518(7540), feb 2015.

Mordatch, I. and Abbeel, P. Emergence of Grounded Compositional Language in Multi-Agent Populations. *CoRR*, abs/1703.0, 2017.

Nayyar, A., Mahajan, A., and Teneketzis, D. Optimal Control Strategies in Delayed Sharing Information Structures. *Automatic Control, IEEE Transactions on*, 56(7), 2011.

Nguyen, D. T., Kumar, A., and Lau, H. C. Policy gradient with value function approximation for collective multi-agent planning. In *NIPS*, pp. 4319–4329. 2017.

Oliehoek, F. A. Sufficient Plan-Time Statistics for Decentralized POMDPs. In *IJCAI*, 2013.

Oliehoek, F. A., Spaan, M. T. J., and Vlassis, N. A. Optimal and Approximate Q-value Functions for Decentralized POMDPs. *Journal of AI Research*, 32, 2008.

Oliehoek, F. A., Spaan, M. T. J., Dibangoye, J. S., and Amato, C. Heuristic search for identical payoff Bayesian games. In *AAMAS*, pp. 1115–1122, 2010.

Oliehoek, F. A., Spaan, M. T. J., Amato, C., and Whiteson, S. Incremental Clustering and Expansion for Faster Optimal Planning in Dec-POMDPs. *Journal of AI Research*, 46, 2013.

Panait, L. and Luke, S. Cooperative multi-agent learning: The state of the art. *Autonomous Agents and Multi-Agent Systems*, 11(3), 2005.

Peshkin, L., Kim, K.-E., Meuleau, N., and Kaelbling, L. P. Learning to Cooperate via Policy Search. In *UAI*, 2000.

Puterman, M. L. *Markov Decision Processes, Discrete Stochastic Dynamic Programming*. Hoboken, New Jersey, 1994.

Radner, R. Team Decision Problems. *Ann. Math. Statist.*, 33(3), 1962.

Robbins, H. and Monro, S. A stochastic approximation method. *The annals of mathematical statistics*, 22(3), 1951.

Rummery, G. A. and Niranjan, M. On-line Q-learning using connectionist systems. Technical report, Cambridge University Engineering Department, 1994.

Salustowicz, R., Wiering, M., and Schmidhuber, J. Learning Team Strategies: Soccer Case Studies. *ML*, 33(2-3), 1998.

Shani, G., Pineau, J., and Kaplow, R. A survey of point-based POMDP solvers. *Journal of Autonomous Agents and Multi-Agent Systems*, 27(1), 2013.

Sutton, R. S. and Barto, A. G. *Introduction to Reinforcement Learning*. Cambridge, MA, USA, 1st edition, 1998. ISBN 0262193981.

Szer, D., Charpillet, F., and Zilberstein, S. MAA*: A Heuristic Search Algorithm for Solving Decentralized POMDPs. In *UAI*, 2005.

Tan, M. Multi-agent Reinforcement Learning: Independent vs. Cooperative Agents. In *Readings in Agents*. San Francisco, CA, USA, 1998.

Watkins, C. J. C. H. and Dayan, P. Q-Learning. *ML*, 8(3), 1992.

Wu, F., Zilberstein, S., and Jennings, N. R. Monte-Carlo Expectation Maximization for Decentralized POMDPs. In *IJCAI*, 2013.

Zhang, C. and Lesser, V. Coordinated Multi-Agent Reinforcement Learning in Networked Distributed POMDPs. In *AAAI*, San Francisco, California, USA, 2011.